

Title: Risk Alert Classifier

Duration: 6 Hours

Objective

The objective of this project is to design, evaluate, and optimize a classification system that predicts high-risk customer behavior. Students will implement multiple classification algorithms, evaluate models using advanced classification metrics, handle imbalanced data, and improve performance using sampling techniques and hyperparameter tuning.

Problem Statement

You are working as a **Data Scientist** for a digital banking platform. The bank wants to build an **early-warning system** that identifies **high-risk customers** who are likely to default on payments or engage in fraudulent behavior.

The dataset contains customer demographic, behavioral, and transaction-related features. However, the dataset is **highly imbalanced**, where risky customers form a small minority.

Your task is to build a **robust classification pipeline** that:

- Accurately identifies high-risk customers
- Handles class imbalance effectively
- Evaluates performance using appropriate classification metrics
- Improves model accuracy using **hyperparameter tuning**

Part A: Conceptual Understanding (Theory)

Answer briefly (2–4 lines each):

1. What is **Logistic Regression** and why is it suitable for classification?
 2. Explain **classification performance metrics** and why accuracy alone is insufficient.
 3. Define **Type-I Error** and **Type-II Error** in the context of risk prediction.
 4. Explain **Precision, Recall, F1-Score, TPR, and FPR**.
 5. What is **AUC-ROC** and how does it help in evaluating classifiers?
 6. Why does **imbalanced data** create problems in classification models?
-

Part B: Dataset Understanding & Preparation

You are provided with a dataset containing:

- Customer demographic information
- Transaction activity indicators
- Credit behavior indicators
- Target variable: **Risk Status**
 - 0 → Low Risk
 - 1 → High Risk

Dataset: [Access from here](#)

Tasks:

7. Identify input features and target variable.
 8. Perform a **train-test split** while maintaining class distribution.
 9. Identify missing values and apply **KNN Imputer** for multivariate imputation.
-

Part C: Baseline Classification Model

10. Implement **Logistic Regression** as a baseline model.
 11. Generate and interpret:
 - Confusion Matrix
 - Accuracy Score
 - Precision, Recall, F1-Score
 12. Identify **Type-I and Type-II errors** from the confusion matrix.
-

Part D: Handling Imbalanced Data

13. Demonstrate the impact of class imbalance on model performance.
 14. Apply the following techniques and retrain the model:
 - Under-Sampling
 - Over-Sampling
 - SMOTE
 - ADASYN
 15. Compare performance before and after balancing using:
 - Recall for minority class
 - F1-Score
 - AUC-ROC
-

Part E: Tree-Based Classification Models

16. Implement **Decision Tree Classifier**.
 17. Analyze overfitting by comparing training and testing performance.
 18. Implement **Random Forest Classifier**.
 19. Compare Decision Tree vs Random Forest in terms of accuracy and generalization.
-

Part F: Hyperparameter Tuning

20. Apply **Randomized Search CV** to optimize:
 - Decision Tree hyperparameters
 - Random Forest hyperparameters
 21. Apply **Grid Search CV** for fine-tuning the best performing model.
 22. Compare tuned vs untuned model performance.
-

Part G: Model Evaluation & ROC Analysis

23. Plot and interpret the **ROC Curve** for all models.
24. Compute and compare **AUC-ROC scores**.
25. Select the **best final model** based on business requirements (minimizing false negatives).

Part H: Final Analysis & Reporting

26. Prepare a final report including:
- Best classification model and justification
 - Impact of imbalance handling techniques
 - Comparison of performance metrics
 - Business interpretation of false positives and false negatives
27. Submit:
- Source code / Jupyter Notebook
 - Evaluation tables and plots
 - Final conclusions and recommendations
-

Submission Guidelines

- Include practical implementation in a Jupyter Notebook and screenshots.
- Label all charts clearly and write short interpretations under each result.
- GitHub Repository:
 - Create a GitHub repository to host your project.
 - Upload your project files, including source code, and documentation to the repository.
 - Add a document (PDF) explaining theory concepts with definitions (if any).
 - Ensure that you provide a clear and descriptive README.md file.

Remember to follow the instructions provided professionally, make suitable assumptions wherever necessary, and avoid copying code or content from unauthorized sources. Good luck with your project work!

Risk Alert Classifier
Supervised Learning

BRING ON YOUR CODING ATTITUDE